

PKU–Zurich PhD Summer School on Machine Learning for Macroeconomics and Finance

ReLoBRaLo and gradient-conflict mitigation in multi-equation residual losses

Simon Scheidegger

HEC Lausanne · University of Zurich

July 6–10, 2026 · Beijing, China

https://liboecon.com/2026_summer_school.html

Roadmap of This Lecture

I. Motivation

- ▶ Multi-component losses in PINNs and DEQNs
- ▶ The scale problem
- ▶ Multi-output architectures

II. Balancing Methods

- ▶ Equal / manual weighting
- ▶ Inverse-loss weighting
- ▶ SoftAdapt
- ▶ **ReLoBRaLo** (+ finger exercise)
- ▶ GradNorm (brief)

III. Experiments & Wrap-up

- ▶ Setup and results
- ▶ Sensitivity to temperature T
- ▶ Connection to PINNs & DSGE
- ▶ Summary and references

Hands-on: 02_03_Loss_Normalization.ipynb

Learning objectives

By the end of this lecture you should be able to:

- ▶ Identify scale problems in multi-component losses (PINNs, DEQNs, multi-task) and diagnose gradient dominance
- ▶ Implement adaptive loss-balancing strategies: inverse-loss weighting, SoftAdapt, and ReLoBRaLo
- ▶ Derive and apply the ReLoBRaLo relative-ratio update rule with baseline mixing and exponential smoothing
- ▶ Compare GradNorm (gradient-norm balancing) with ReLoBRaLo and select appropriate methods for constrained optimization
- ▶ Tune temperature parameters and validate on multi-scale function approximation; apply to DEQN training with PINN-style residual losses

Part I

Motivation

Motivation: Multi-Component Losses

Many problems in economics require minimizing *several* loss terms simultaneously:

Physics-informed NNs (PINNs):

- ▶ PDE residual $\mathcal{L}_{\text{PDE}}$
- ▶ Boundary conditions $\mathcal{L}_{\text{BC}}$
- ▶ Initial conditions $\mathcal{L}_{\text{IC}}$

Deep equilibrium nets (DEQNs):

- ▶ Euler equation residuals
- ▶ Budget constraints
- ▶ Market clearing

Multi-task learning:

- ▶ Predict multiple outputs from shared features
- ▶ Each output has its own loss

Core challenge

These loss components often live at **vastly different scales**.

The Scale Problem

Example: Three loss components


$\mathcal{L}_A \approx 10^{-2}$ (small-scale function)

$\mathcal{L}_B \approx 10^2$ (medium-scale)


$\mathcal{L}_C \approx 10^4$ (large-scale)

With equal weights: $\mathcal{L}_C$ contributes $10^6 \times$ more gradient signal than $\mathcal{L}_A$.

$\Rightarrow$ The network *ignores* the small-scale component.



Multi-Output Network Architecture



$$\mathcal{L}_{\text{total}} = w_A \cdot \mathcal{L}_A + w_B \cdot \mathcal{L}_B + w_C \cdot \mathcal{L}_C$$

How should we set the weights w_A, w_B, w_C ?

Loss Balancing Strategies

Equal Weighting

Simplest approach: $w_A = w_B = w_C = 1$

$$\mathcal{L} = \mathcal{L}_A + \mathcal{L}_B + \mathcal{L}_C$$

Problem:

- ▶ If $\mathcal{L}_C \gg \mathcal{L}_A$, gradients are **dominated** by $\mathcal{L}_C$
- ▶ Network learns to minimize $\mathcal{L}_C$ first
- ▶ $\mathcal{L}_A$ barely decreases
- ▶ Small-scale structure is *ignored*


Gradient dominance

$$\frac{\partial \mathcal{L}}{\partial \theta} \approx \frac{\partial \mathcal{L}_C}{\partial \theta} \quad \text{when } \mathcal{L}_C \gg \mathcal{L}_A, \mathcal{L}_B.$$

The network effectively solves a *single-task* problem.

Equal Weighting — Training Dynamics

Equal Weighting



Generated by 02_03_Loss_Normalization.ipynb (Section 3, equal-weight training). Left: per-component MSE on a log axis; right: the (constant) component weights. $\mathcal{L}_C$ (scale 10^4) drives the gradient and converges quickly; $\mathcal{L}_B$ (scale 10^2) follows; $\mathcal{L}_A$ (scale 10^0) is **essentially not learned**, exactly as the analytical scale argument predicts.

Manual Weighting

Idea: Choose weights to approximately equalize the loss magnitudes:

$$\mathcal{L} = w_A \cdot \mathcal{L}_A + w_B \cdot \mathcal{L}_B + w_C \cdot \mathcal{L}_C$$

e.g., $w_A = 10^4$, $w_B = 10^0$, $w_C = 10^{-2}$.

Pros:

- ▶ Can work well if scales are known a priori
- ▶ Simple to implement

Cons:

- ▶ **Trial-and-error**: need to tune d extra hyperparameters
- ▶ **Fragile**: optimal weights change during training as losses evolve
- ▶ **Doesn't generalize**: weights are problem-specific

Use adaptive methods that adjust weights during training.

Inverse-Loss Weighting

Idea: Weight each component inversely proportional to its current magnitude:

$$w_i^{(t)} = \frac{1}{\mathcal{L}_i^{(t)}}$$

This automatically up-weights small losses and down-weights large ones.

With exponential averaging (for stability):

$$\bar{\mathcal{L}}_i^{(t)} = \beta_{\text{ema}} \cdot \bar{\mathcal{L}}_i^{(t-1)} + (1 - \beta_{\text{ema}}) \cdot \mathcal{L}_i^{(t)}$$

$$w_i^{(t)} = \frac{1}{\bar{\mathcal{L}}_i^{(t)} + \epsilon}$$

Pros:

- ▶ Simple, no hyperparameters (besides β_{ema})
- ▶ Adaptive to changing scales

Cons:

- ▶ **Unstable** when $\mathcal{L}_i \rightarrow 0$
- ▶ Weights can **oscillate** wildly
- ▶ No notion of *relative progress*

Inverse-Loss Weighting — Summary

Algorithm

At each epoch t :

1. Compute current losses $\mathcal{L}_1^{(t)}, \dots, \mathcal{L}_K^{(t)}$
2. Update exponential averages: $\bar{\mathcal{L}}_i^{(t)} = \beta_{\text{ema}} \bar{\mathcal{L}}_i^{(t-1)} + (1 - \beta_{\text{ema}}) \mathcal{L}_i^{(t)}$
3. Set weights: $w_i^{(t)} = \frac{1}{\bar{\mathcal{L}}_i^{(t)} + \epsilon}$
4. Normalize: $\tilde{w}_i^{(t)} = K \cdot \frac{w_i^{(t)}}{\sum_{j=1}^K w_j^{(t)}}$

The normalization ensures $\sum_i \tilde{w}_i = K$ (total weight is preserved).

Typical choices: $\beta_{\text{ema}} = 0.99$, $\epsilon = 10^{-8}$.

SoftAdapt

Idea: Weight based on *relative loss changes* — components that are improving slowly get higher weight.

Rate of change

$$\Delta_i^{(t)} = \mathcal{L}_i^{(t)} - \mathcal{L}_i^{(t-1)}$$

- ▶ $\Delta_i > 0$: loss *increased* $\Rightarrow$ needs more weight
- ▶ $\Delta_i < 0$: loss *decreased* $\Rightarrow$ already learning well

Softmax temperature T :

- ▶ Small T : winner-take-all
- ▶ Large T : nearly equal weights

Pros:

- ▶ Focuses on *struggling* components
- ▶ Simple to implement

Cons:

- ▶ Sensitive to T
- ▶ Noisy loss changes $\Rightarrow$ noisy weights

SoftAdapt — Formula

Weight update (Heydari et al., 2019)

$$\hat{w}_i^{(t)} = \frac{\exp(\Delta_i^{(t)} / T)}{\sum_{j=1}^K \exp(\Delta_j^{(t)} / T)} \cdot K$$

where $\Delta_i^{(t)} = \mathcal{L}_i^{(t)} - \mathcal{L}_i^{(t-1)}$ and $T > 0$ is a temperature parameter.

Interpretation:

- ▶ The softmax converts relative loss changes into a probability-like weighting
- ▶ Multiplication by K ensures $\sum_i \hat{w}_i = K$ (preserves total scale)
- ▶ Components with **increasing loss** (poor progress) receive more weight
- ▶ Components with decreasing loss receive less weight

Scale caveat

Raw $\Delta_i^{(t)}$ is dimensional: a component with loss $\sim 10^3$ produces fluctuations that swamp one with loss $\sim 10^{-3}$. In practice rescale before the softmax, e.g. $\tilde{\Delta}_i^{(t)} = \Delta_i^{(t)} / (\mathcal{L}_i^{(t)} + \epsilon)$, so the mechanism reacts to *relative* progress rather than absolute magnitudes.

ReLoBRaLo — Relative Loss Balancing with Random Lookback

Key innovations (Bischof & Kraus, 2025), adapted here as a deterministic classroom version:

1. **Relative balancing:** Use *ratios* $\mathcal{L}_i^{(t)} / \mathcal{L}_i^{(t-1)}$ instead of absolute values
2. **Baseline comparison:** Mix in the **initial losses** $\mathcal{L}_i^{(0)}$ as a stable reference
3. **Exponential smoothing:** Blend new weights with history via α

Why random lookback?

- ▶ Prevents weight drift
- ▶ “Resets” to global perspective
- ▶ Stabilizes training
- ▶ Stochastic regularization of the baseline-vs-step trade-off

ReLoBRaLo — Algorithm

Weight update at epoch t

- 1: Compute losses $\mathcal{L}_1^{(t)}, \dots, \mathcal{L}_K^{(t)}$
- 2: **Step-wise ratios:** $r_{i,s}^{(t)} = \mathcal{L}_i^{(t)} / (T \mathcal{L}_i^{(t-1)} + \varepsilon)$
- 3: **Baseline ratios:** $r_{i,b}^{(t)} = \mathcal{L}_i^{(t)} / (T \mathcal{L}_i^{(0)} + \varepsilon)$
- 4: $\hat{w}_{i,s}^{(t)} = K \text{ softmax}_i(r_{\cdot,s}^{(t)})$, $\hat{w}_{i,b}^{(t)} = K \text{ softmax}_i(r_{\cdot,b}^{(t)})$
- 5: **Combine with smoothing:**

$$w_i^{(t)} = \alpha \left[\rho w_i^{(t-1)} + (1 - \rho) \hat{w}_{i,b}^{(t)} \right] + (1 - \alpha) \hat{w}_{i,s}^{(t)}$$

- ▶ α : smoothing — how much to trust the historical weights
- ▶ ρ : baseline mix — relative weight on the initial-loss comparison
- ▶ T : temperature — sharpness of the softmax

Finger Exercise: Loss Balancing

Exercise

A DEQN has three loss components at iteration t : $\mathcal{L}_1^{(t)} = 10.0$, $\mathcal{L}_2^{(t)} = 0.01$, $\mathcal{L}_3^{(t)} = 1.0$.

1. If you use **equal weights** $w_i = 1/3$, which loss component dominates the gradient? Is this desirable?
2. Compute the **inverse-loss weights**: $w_i \propto 1/\mathcal{L}_i^{(t)}$, normalized to sum to 1.
3. Which component now receives the highest weight? Why does this help training?

Solution: Loss Balancing

Answer

1. With equal weights, $\mathcal{L}_1 = 10.0$ dominates ($100\times$ larger than $\mathcal{L}_2$). The optimizer focuses almost entirely on reducing $\mathcal{L}_1$ while ignoring $\mathcal{L}_2$. This is undesirable if all three conditions must be satisfied simultaneously.
2. Unnormalized: $1/10.0 = 0.1$, $1/0.01 = 100$, $1/1.0 = 1$. Sum = 101.1.
Normalized: $w_1 = 0.001$, $w_2 = 0.989$, $w_3 = 0.010$.
3. $\mathcal{L}_2$ receives the highest weight ($\approx 99\%$). Inverse-loss weighting amplifies the components that are *hardest to satisfy* (smallest loss = most satisfied $\Rightarrow$ smallest weight), directing optimization effort where it is most needed.

ReLoBRaLo — Hyperparameters

Symbol	Name	Default	Effect
T	Temperature	1.0	Softmax sharpness: small $T \rightarrow$ winner-take-all; large $T \rightarrow$ uniform weights.
α	Smoothing	0.999	Trust in historical weights; larger $\alpha \rightarrow$ slower adaptation.
ρ	Baseline mix	0.999	Weight on initial-loss comparison (convex blend).


Note. The original algorithm uses a stochastic Bernoulli lookback of probability ρ ; the classroom version is a deterministic convex blend (simpler to debug). T , α , ρ are *balancing hyperparameters* from the ReLoBRaLo literature — not Session 1's economic symbols (α capital share, β discount factor); likewise $\mathcal{L}_i$ is a loss *component*, Session 1's scalar loss ℓ_θ being the single-component case.

Practical guidance

Defaults work well in most cases. Only T needs occasional tuning (see sensitivity analysis in Notebook 02_03).

ReLoBRaLo — Weight Evolution

ReLoBRaLo ($T=1.0$)



Generated by 02_03_Loss_Normalization.ipynb (Section 5, ReLoBRaLo at $T = 1$). Left: per-component MSE on a log axis under ReLoBRaLo (compare with the equal-weighting plot on slide 7). Right: the actual weight trajectories w_A, w_B, w_C produced by the ReLoBRaLo update rule. w_A rises (small-scale loss now amplified), w_C falls (large-scale loss attenuated), exactly as the algorithm targets.

GradNorm (Brief)

Idea (Chen et al., 2018): Balance the *gradient norms* across tasks, not the losses.

Objective

Adjust weights w_i so that:

$$\|\nabla_{\theta} w_i \mathcal{L}_i\|_2 \approx \bar{G} \cdot r_i^{(t)}$$

where $\bar{G}$ is the average gradient norm and $r_i^{(t)}$ captures relative training rates.

The weights w_i are **learnable parameters** updated via gradient descent on a “gradient balance” loss.

Pros:

- ▶ Principled — directly addresses gradient competition
- ▶ Automatic, no manual tuning

Cons:

- ▶ Requires per-task gradient computation (expensive)
- ▶ More complex to implement
- ▶ Hyperparameter α (asymmetry)

GradNorm — Formula

Gradient normalization loss (Chen et al., 2018)

$$\mathcal{L}_{\text{grad}} = \sum_{i=1}^K \left| \|\nabla_{\theta} w_i \mathcal{L}_i\|_2 - \bar{G} \cdot r_i^{(t)} \right|$$

where:

- ▶ $\bar{G} = \frac{1}{K} \sum_i \|\nabla_{\theta} w_i \mathcal{L}_i\|_2$ — average gradient norm
- ▶ $r_i^{(t)} = (\tilde{\mathcal{L}}_i^{(t)} / \tilde{\mathcal{L}}^{(t)})^{\alpha}$ — relative inverse training rate
- ▶ $\tilde{\mathcal{L}}_i^{(t)} = \mathcal{L}_i^{(t)} / \mathcal{L}_i^{(0)}$ — loss ratio relative to initial
- ▶ α controls asymmetry ($\alpha = 1.5$ typical)

Update: $w_i \leftarrow w_i - \eta_w \cdot \nabla_{w_i} \mathcal{L}_{\text{grad}}$, then renormalize to $\sum_i w_i = K$.

Experiment: Multi-Scale Function Approximation

Experiment Setup

Three target functions on $[0, 1]^2$:

$$f_A(\mathbf{x}) = \text{genz_gaussian}(\mathbf{x}) \\ \text{scale} \in [0, 1]$$

$$f_B(\mathbf{x}) = 100 \times \text{genz_oscillatory}(\mathbf{x}) \\ \text{scale} \in [-100, 100]$$

$$f_C(\mathbf{x}) = 10,000 \times \text{genz_continuous}(\mathbf{x}) \\ \text{scale} \in [0, 10,000]$$

Architecture:

- ▶ Shared trunk: 3×64 ReLU
- ▶ 3 output heads (Dense(1) each)
- ▶ Adam, $\eta = 10^{-3}$

Data:

- ▶ $n_{\text{train}} = 1,000$
- ▶ $n_{\text{test}} = 2,000$
- ▶ $d = 2$

$$\mathcal{L} = w_A \cdot \text{MSE}(f_A, \hat{f}_A) + w_B \cdot \text{MSE}(f_B, \hat{f}_B) + w_C \cdot \text{MSE}(f_C, \hat{f}_C)$$


Results: Equal vs. Inverse-Loss vs. ReLoBRaLo

	Equal weights	Inverse-loss	ReLoBRaLo
f_A mean error	3.3×10^{-1}	7.1×10^{-3}	1.6×10^{-1}
f_B mean error	5.4×10^1	3.4×10^1	3.1×10^1
f_C mean error	1.2×10^3	4.1×10^3	6.3×10^2
Total rel. error $\sum_i \overline{e_i} /\text{scale}_i$	0.99	0.76	0.53

- **Equal:** f_C dominates the gradient; f_A is the **least accurate** run and the total error is highest.
- **Inverse-loss:** sharply improves f_A and roughly halves the total error, but f_C gets noticeably worse.
- **ReLoBRaLo:** best on f_B and f_C and the **lowest total relative error** ($\sim 1.9\times$ below equal weighting); f_A is fit well, though less aggressively than inverse-loss.

Numbers from 02_03_Loss_Normalization.ipynb, RUN_MODE = "teaching" (500 epochs), seed 42. Absolute values are seed-/mode-sensitive; the qualitative point is that adaptive weighting lowers the scale-weighted total.

Sensitivity to Temperature T



Generated by 02_03_Loss_Normalization.ipynb (Section 7, temperature sweep). Bars show the total relative error $\sum_i \text{MAE}_i / \text{scale}_i$ at the end of training, swept over $T \in \{0.01, 0.1, 1.0, 10.0\}$.

$T \approx 1$ is a robust default; sensitivity analysis is cheap.

T too small ($\ll 1$):

- ▶ Winner-take-all
- ▶ Only one component trained at a time
- ▶ Slow convergence

T too large ($\gg 1$):

- ▶ Near-equal weights
- ▶ Reverts to the scale problem
- ▶ No effective balancing

Connection to PINNs and DSGE

Physics-informed NNs:

- ▶ PDE residual: $\mathcal{L}_{\text{PDE}} \sim \mathcal{O}(1)$
- ▶ Boundary conditions: $\mathcal{L}_{\text{BC}} \sim \mathcal{O}(10^{-2})$
- ▶ Initial conditions: $\mathcal{L}_{\text{IC}} \sim \mathcal{O}(10^{-4})$
- ▶ ReLoBRaLo ensures BCs/ICs are **not ignored**

DSGE / Deep equilibrium nets:

- ▶ Euler equations at different scales
- ▶ Budget constraint residuals
- ▶ Market clearing conditions
- ▶ Each may have different units and magnitudes

Bischof & Kraus (2025)

ReLoBRaLo was originally developed for PDE-constrained optimization and has been shown to improve convergence by up to an order of magnitude on multi-physics problems.

Summary

Key takeaways:

1. Multi-component losses at different scales cause **gradient dominance**.
2. Equal weighting fails when scales differ by $> 10\times$.
3. Inverse-loss weighting helps but can be unstable.
4. **ReLoBRaLo** provides robust, adaptive balancing via relative ratios + baseline mix.
5. Critical for PINNs, DEQNs, and multi-task learning in economics.

References:

- ▶ Bischof & Kraus (2025). *Multi-Objective Loss Balancing for Physics-Informed Deep Learning*, CMAME 439, 117914.
- ▶ Chen et al. (2018). *GradNorm*, ICML.
- ▶ Heydari et al. (2019). *SoftAdapt*, arXiv:1912.12355.

Hands-on

⇒ `02_03_Loss_Normalization.ipynb`

Applied exercise: `code/03_irbc/03_03_IRBC_Exercise.ipynb` uses inverse-loss weighting on the IRBC model of Session 3 — a direct preview of ReLoBRaLo in a DEQN.